



0. Deep Learning

Deep Learning - Why does it renovate Artificial Intelligence?

Contents:

- [1. Deep Learning works with Unstructured Data](#)
- [2. Deep Learning extracts features on its own](#)
- [3. The "Heart" of Deep Learning is Neural Networks](#)

[Summary](#)

[References:](#)

In this section, we're going to talk about Deep Learning, and don't worry—it sounds intimidating, but it really comes from a very simple problem. Deep Learning exists because Machine Learning wasn't enough for the kind of data we deal with in the real world.

1. Deep Learning works with Unstructured Data

When you first learn Machine Learning, everything feels clean and organized: datasets come in tables, with rows and columns, like something you could open in Excel.

- Each column is a feature, each row is an example, and models learn patterns from that structure.

This works beautifully because Machine Learning relies on structure—it expects everything to already be neatly defined.

But reality doesn't look like that at all.

When you use your phone or browse the internet, most of the data you interact with is not tables—it's images, videos, text, and audio. These are what we call unstructured data, and they don't come with obvious columns like "feature 1" or "feature 2".

- An image, for example, is just a grid of pixel values. To a human, it's instantly recognizable—a cat, a car, a face—but to a traditional Machine Learning model, it's just numbers with no meaning attached.

That's where the problem begins: Machine Learning doesn't naturally understand structure in this kind of data, so it depends on humans to manually define features like edges, textures, or shapes.


In other words, before the model even starts learning, you have to tell it what to look for, and that becomes incredibly difficult as data grows more complex.

This is exactly where Deep Learning comes in.

Deep Learning is a subset of Machine Learning designed to handle complex, high-dimensional, and non-linear data—meaning data that does not come in a clear or predefined structure.

Instead of relying on humans to manually define features, Deep Learning models learn these features automatically from raw data through multiple layers of transformation.

In other words, rather than being told what patterns to look for, the model discovers them on its own, making it far more effective for tasks involving images, text, audio, and other messy real-world data where structure is not obvious.



Now add another layer to the problem: **scale**.

As datasets become enormous—millions or even billions of data points—manual feature engineering becomes not just hard, but practically impossible.

Traditional Machine Learning models also tend to rely on relatively simple relationships, often leaning toward linear or semi-linear patterns. That means when the underlying patterns in the data are highly complex (like recognizing objects in images or understanding human language), these models start to struggle.

So the real question becomes:

- What if the machine didn't need us to define the structure at all?
- What if it could learn the structure by itself, directly from raw data?

This is exactly where Deep Learning comes in. Instead of separating the process into “*human extracts features*” and “*model learns patterns*,” Deep Learning combines everything into one system. You can think of the difference like this:

Machine Learning: **Human** → **extract features** → **ML model** → **prediction**

Deep Learning: **Raw data** → **Deep Learning model** → **prediction**

The key idea is that Deep Learning models automatically discover the important features during training. They don't need you to tell them what edges or shapes to look for—they figure that out on their own. And the reason they can do this is because they are built using multiple layers, which is why the term “deep” is used. Each layer learns a different level of abstraction:

- Early layers might detect simple patterns like edges;
- Middle layers combine those into shapes;
- Deeper layers recognize complex objects.

It's like building understanding step by step, from raw pixels all the way to meaningful concepts. A simple way to visualize this process is: **Image** → **edges** → **shapes** → **parts** → **object**

This layered learning process is very similar to how humans learn:

- You don't define strict rules for recognizing a cat—you just see many examples, and your brain gradually learns the patterns.
- Deep Learning models try to mimic this behavior by learning directly from data, especially when that data is large and unstructured.
- So in the end, Deep Learning is not something completely separate from Machine Learning—it's more like an evolution. Machine Learning works well when data is structured and manageable, but Deep Learning takes over when things become too complex, too large, or too unstructured for humans to handle manually.

2. Deep Learning extracts features on its own

To understand this more clearly, let's step back and ask a very simple question: What do computers actually see? Because the way we see the world and the way machines see the world are completely different.

Imagine you have two types of data:

- A table of students with their home location stored as coordinates;
- An image of a cat.


To a human, this is straightforward.

- You look at the table and think, *“Okay, these are students, and those numbers represent locations... though I might need the dataset creator to explain exactly what those coordinates mean”*.
- And when you see the image, you immediately say, “That's a cat,” because you've seen cats before and your brain already understands the concept.

But to a computer, both of these look equally meaningless at first.

- The table is just “a bunch of columns with numbers,” and the image is just “a matrix of pixel values.”
- There is no built-in understanding of what a “student,” a “location,” or a “cat” is. It doesn't see meaning—it only sees data.

If you go even lower level, a document or text file is just a long sequence of bits like “1010101010”.



WHAT COMPUTERS SEE VS WHAT HUMANS SEE

So the real challenge in AI is not just learning patterns—it's creating meaning from raw data.

Now, this is where Machine Learning and Deep Learning start to diverge in philosophy. In traditional Machine Learning, you usually help the model by structuring the data in a meaningful way.

- For example, if you have student coordinates, you might use an algorithm like K-Means clustering, and it can learn to group students into regions like West, East, or South-East based on those numeric features:
- That works because the data is already clean and numerical.
- But notice something important: The model didn't "understand" locations—it just found patterns in numbers that you provided in a structured format.

We can summarize that idea like this:

Human defines features → ML finds patterns → output (clusters, predictions)

Now let's move to images. Suppose you give a Machine Learning model a cat image:

- It doesn't know what a cat is, so you would traditionally need to extract features manually—edges, textures, shapes, color histograms.
- Again, the human is doing the hard work of deciding what matters. The model just learns relationships between those features.

But Deep Learning flips this idea completely. Instead of you telling the model what features to use, the model learns them by itself. When you give it an image:

- It starts from raw pixels and gradually builds understanding through layers.
- Early layers might detect simple things like edges or color gradients, then later layers combine those into shapes, and eventually into meaningful parts like "this region might be an eye" or "this looks like ears."

- So instead of you saying “this is an eye,” the model figures it out on its own.

You can think of it like this:

Raw image → pixel patterns → edges → shapes → parts → “cat”

Now let’s go back to structured vs unstructured data with another example. If your dataset contains coordinates, Machine Learning works fine. But what if instead of coordinates, you store full addresses like “2 Saint Louis Street”?

Now things get messy. That’s text, not numbers. A traditional ML model will struggle unless you manually convert that text into numerical features—maybe encoding regions, extracting city names, or mapping addresses to coordinates yourself.

In other words, you are forced to reshape the data into something the model can understand. Deep Learning, however, can work directly with that raw text. It can learn patterns in language, extract meaning, and even understand relationships between words—without you explicitly defining every feature.

- It might learn that “Street,” “Avenue,” and “Road” are related concepts, or that certain address patterns correspond to certain regions.

So this brings us to the most important idea, and I really want to emphasize this:

In Machine Learning, features are created by humans. In Deep Learning, features are learned by the model itself.

That’s the fundamental shift. Machine Learning is like giving someone a fully prepared meal and asking them to analyze the taste. Deep Learning is like giving raw ingredients and letting the system figure out not only how to cook, but also what flavors matter.

And the reason this works—especially for images, text, and other unstructured data—is because Deep Learning models use multiple layers to progressively transform raw data into meaningful representations.

So when you hear “Deep Learning,” don’t think of it as something magical. It’s simply a system that says: “Give me the raw data... I’ll figure out what’s important.”

3. The “Heart” of Deep Learning is Neural Networks

So now that we understand the key idea—Machine Learning depends on humans to define structure, while Deep Learning learns that structure by itself—we can finally dive into the core engine that makes this possible: **Neural Networks**

And yes, they are inspired by the human brain, but don’t overthink that comparison. You don’t need biology here—just think of it as a system that takes input, processes it through layers, and produces an output.

At the most basic level, a Neural Network follows a simple structure:

Input → HiddenLayers → Output

Let’s go back to our cat example, because it makes everything much clearer. Imagine you feed

an image of a cat into the network:

- The input layer receives raw pixel values—just numbers, nothing meaningful yet.
- Then the data flows through multiple hidden layers, where the real learning happens
 - Each connection between neurons has something called a weight, which basically answers the question: “Is this feature important?”
 - If a certain pattern (like a curve that might be an ear) is useful, its weight increases.
 - If it's not useful, the weight decreases.— Over time, the network learns which patterns matter and which don't.
- Now, after passing through all these layers, the network makes a prediction at the output layer.
- Maybe it says: “This is a cat with 0.8 probability, a dog with 0.2 probability.”

This is where **activation functions** come in—things like Sigmoid or Softmax help convert raw outputs into probabilities, because at the end of the day, the model is always making a guess with some level of confidence.

- Of course, the first guess is usually wrong. That's where the loss function comes in. It measures how far the prediction is from the actual answer.
 - If the image really is a cat but the model isn't confident, the loss will be high.
- And just like in Machine Learning, the goal is simple: minimize the error. To do that, we use an optimizer, which updates the weights step by step, nudging the model toward better predictions.

If you want to see the whole workflow clearly, it looks like this:

Input (cat image) → Predict → Compute Loss → Optimizer updates weights → Repeat

And this loop happens again and again, thousands or even millions of times, until the model becomes very good at recognizing cats.

The beauty here is that at no point did we explicitly tell the model what “ears” or “eyes” are—it discovered those features by itself through layers and weight adjustments.

Summary

So when we connect everything together, the idea becomes very clean:

Machine Learning depends on humans to define structure, while Deep Learning learns that structure by itself. And that single difference is what makes it powerful enough to handle images, videos, language, and all the messy data that actually exists in the real world. And Neural Networks make it possible by providing the mechanism for this automatic learning—layers that transform raw data into meaningful representations, weights that measure importance, activation functions that interpret predictions, and optimization processes that continuously improve the model. Without Neural Networks, Deep Learning would just be an idea with no way to execute.

In the section closely related to this one, we'll go deeper into Neural Networks themselves—how they are built, how they actually learn step by step, and why they are so powerful despite being conceptually simple.

References:

<https://www.geeksforgeeks.org/deep-learning/introduction-deep-learning/>

<https://www.ibm.com/think/topics/deep-learning>

Next Section:

 [1. Neural Networks and Perceptrons](#)