



28. Support Vector Machines

Making decisions with a hyperplane?

Previous Section:

 [27. Independent Component Analysis](#)

[1. Support Vector Machines: The Overview](#)

[2. Support Vector Machines: The Workflow](#)

[3. Support Vector Machines with Python](#)

[4. Support Vector Machines: Kernel Methods](#)

[5. Margins in Support Vector Machine](#)

[Summary](#)

[References](#)

This section is about Support Vector Machines, one of the most powerful models in Machine Learning—and also a quite challenging one to decipher. Especially when it comes to what the model can actually do.

At first glance, it sounds simple. ***“Separate the data with a hyperplane.”*** → That’s it? Hardly that simple. Because just like many things in Machine Learning, the simplicity is only on the surface.

You might hear people say SVM is “powerful,” “versatile,” or even “elegant”. But let’s be honest for a moment—those words don’t really tell you anything. Every model gets described that way at some point. So instead of throwing adjectives around, let’s slow down and actually understand what’s happening.

At its core, Support Vector Machines are about making decisions. Not by memorizing data. Not by building complex trees. But by drawing a boundary—a line, or in higher dimensions, a hyperplane—that separates different classes.

Now here’s the interesting part. You’ve already seen something like this before. Think back to how some models create decision boundaries:

- A line separating two groups of points
- A curve dividing regions
- A rule that says **“this side = class A, that side = class B”**

SVM does something very similar, but with a twist. It doesn't just find *any* boundary. It finds the **best possible boundary**. And “best” here has a very specific meaning—one that leads us directly into the two ideas that define SVM:

- **Margins** → how far the boundary is from the closest data points
- **Kernel methods** → how we handle data that cannot be separated in a simple way

So if the phrase “hyperplane” sounds intimidating, don’t worry. By the end of this section, you’ll see it for what it really is: **Just a decision boundary drawn with a very careful strategy.**


1. Support Vector Machines: The Overview

Support Vector Machines (SVM) is a supervised machine learning algorithm used for classification and regression tasks. Its core objective is to find the best possible boundary—known as a **hyperplane**—that separates different classes in the dataset.

But here’s the important shift in thinking. SVM does **not** just try to separate the data. It tries to separate the data in the **most optimal way possible**.

Instead of minimizing classification errors directly, SVM focuses on **maximizing the separation between classes**. And that single idea—maximize the distance between groups—is what gives SVM its strong ability to generalize to unseen data.

(The image below is for reference only)



The Core Components

Let’s break down what actually makes SVM work.

- **Hyperplane**

A hyperplane is the decision boundary that separates different classes in the feature space. In the linear case, it is defined as:

$$w \cdot x + b = 0$$

Where: w determines the **orientation** of the boundary; b shifts the boundary in space

In 2D, it looks like a line. In 3D, it becomes a plane. In higher dimensions, it's still the same idea—just harder to visualize.

- **Support Vectors**

These are the **most important data points in the entire model**. They are the points **closest to the hyperplane**, and they directly determine where the boundary is placed.

Here's the surprising part: Remove other data points, the boundary stays the same. Remove support vectors, the boundary changes completely

So in a way, SVM doesn't care about all data equally—it only "listens" to the most critical ones.

- **Margin**

The margin is the distance between the hyperplane and the closest data points (the support vectors). SVM's true objective is: **Maximize this margin**. Why? Because a larger margin means better separation, more robustness, and lower risk of overfitting

- **Kernel**

What if the data is not linearly separable? That's where kernels come in. A kernel function **implicitly transforms data into a higher-dimensional space**, where separation becomes possible.

Common kernels include: Linear, Polynomial, and Radial Basis Function (RBF)

The beauty? You don't actually compute the transformation—you just **pretend you did**. (This is the famous *kernel trick*.)

- **Hard Margin**

This is the ideal scenario: Perfect separation → No misclassification allowed

But in reality? Almost useless. Because real-world data is noisy, messy, and overlapping.

- **Soft Margin**

This is the practical version of SVM. Instead of forcing perfect separation, it allows some misclassification → some points inside the margin. The goal becomes a balance → **Wide margin + acceptable errors**

- **C (Regularization Parameter)**

This parameter controls the trade-off:

Large C → strict, fewer errors, smaller margin

Small C → relaxed, more errors, wider margin

So C basically answers → *"Do we care more about correctness... or generalization?"*

- **Hinge Loss**

SVM uses a special loss function that penalizes misclassified points and penalizes points inside the margin. It doesn't just care if you're wrong—it cares **how wrong you are relative to the margin**.

- **Dual Problem**

Instead of solving the problem directly, SVM reformulates it using **Lagrange multipliers**. Why? Because this allows efficient computation, the use of kernels, and focus only on support vectors

There are two types of SVM

- **Linear SVM:** Uses a straight-line (or flat) boundary → Fast and efficient → Works well when data is roughly separable
- **Non-Linear SVM:** Uses kernels to create complex boundaries → Can handle highly structured, non-linear data

If I had to choose the most important component in SVM—besides the hyperplane—it would be the **support vectors**. Because they are the ones that *define everything*.

For each class, only a few points actually matter. And SVM's entire goal is to position the hyperplane such that:

*The distance between the hyperplane and these support vectors is **as large as possible**. Not small. Not average. **Maximized**.*

And one last thing. A hyperplane is not just “a line.” → Think of it more like a **geometric structure**. In 1D, it definitely looks like a line. In 2D, it is a plane. But conceptually, it behaves more like a **boundary region**, almost like drawing a tight rectangle or corridor between two classes.

The hyperplane sits in the middle. And the support vectors define its edges. So when you think of SVM, don't think: “Find a line.” → Think: **Find the safest possible boundary between two worlds**.

2. Support Vector Machines: The Workflow

Now let's stop talking in abstractions for a moment. You've heard about hyperplanes, margins, and support vectors. But what does SVM actually **do** when we give it data? Let's walk through it—step by step.

First, consider the following dataset. We are trying to answer a very simple question: *Should a customer be granted credit?* (Yes / No). But as always, simple questions rarely come with simple data.

Original Dataset (Categorical Form)

ID	Age	Has Job	Own House	Credit Rating	Class
1	Young	false	false	fair	No
2	Young	false	false	good	No
3	Young	true	false	good	Yes
4	Young	true	true	fair	Yes
5	Young	false	false	fair	No
6	middle	false	false	fair	No

ID	Age	Has Job	Own House	Credit Rating	Class
7	middle	false	false	good	No
8	middle	true	true	good	Yes
9	middle	false	true	excellent	Yes
10	middle	false	true	excellent	Yes
11	old	false	true	excellent	Yes
12	old	false	true	good	Yes
13	old	true	false	good	Yes
14	old	true	false	excellent	Yes
15	old	false	false	fair	No

(The ID is just an index. It means nothing to the model.)

Now here comes the first problem. SVM does not understand words like: "**Young**"; "**Good**"; or "**Excellent**" → To SVM, these are just meaningless symbols. It only understands **numbers**.

Converting Categorical Features to Numerical Values - We translate everything into numbers. Not because we want to. But because the model demands it.

Encoded Dataset (Numerical Form)

ID	Age	Job	House	Credit	y
1	0	0	0	0	-1
2	0	0	0	1	-1
3	0	1	0	1	+1
4	0	1	1	0	+1
5	0	0	0	0	-1
6	1	0	0	0	-1
7	1	0	0	1	-1
8	1	1	1	1	+1
9	1	0	1	2	+1
10	1	0	1	2	+1
11	2	0	1	2	+1
12	2	0	1	1	+1
13	2	1	0	1	+1
14	2	1	0	2	+1
15	2	0	0	0	-1

Where:

- **Age: Young = 0, Middle = 1, Old = 2**
- **Job / House: 0 = No, 1 = Yes**
- **Credit: Fair = 0, Good = 1, Excellent = 2**

- **Class:** *Yes* → **+1**, *No* → **-1**

Why +1 and -1? This is not just a random choice. SVM doesn't think in "**Yes**" or "**No**". It thinks in **signs**. And everything revolves around one constraint: $y_i(w \cdot x_i + b) \geq 1$

This condition ensures correct classification AND that each point lies **outside the margin**. If:

- $y_i = +1$ → point must be on one side
- $y_i = -1$ → point must be on the other

So the label isn't just a label. It tells the model **which side of reality the point belongs to**

Decision Function

Once training is done, SVM makes decisions using:

$$f(x) = w \cdot x + b$$

And the prediction is simply:

$$y = \text{sign}(f(x))$$

Positive → **+1**(*Yes*); Negative → **-1**(*No*)

Interpreting the Components

Now let's slow down and actually understand what these symbols mean.

- **Feature Vector** x (generally $x = (x_1, x_2, \dots, x_d)$)

This represents one customer: $x = (\text{Age}, \text{Job}, \text{House}, \text{Credit})$

Each value is now numerical. To us, it's still a person. To SVM, it's just a point in space.

- **Weight Vector** w (generally $w = (w_1, w_2, \dots, w_d)$)

This is what the model learns. Each weight answers a silent question: *How important is this feature?*

- Large $|w_j|$ → strong influence
- Small $|w_j|$ → weak influence

So SVM is not just separating data. It is **deciding what matters**.

- **Bias** b

If w defines direction. Then b defines position. It shifts the hyperplane: Left / Right or Up / Down, or in higher dimensions through space

So what is SVM actually doing here? It takes each customer. Turns them into a point. And asks:

┃ *Which side of the boundary do you belong to?*

But not just any boundary. A boundary that maximizes separation, respects the most critical points, and prepares for data it has never seen before

And that's the workflow. Not complicated. But very precise. Now, the real question begins: *What does this data actually tell us?*

Interpreting the Training Data

Now this is where things start to get interesting. We're no longer just preparing data. We're starting to **understand what the model is actually seeing**.

In this example, each customer is represented as a **feature vector** x_i ; a **label** $y \in \{-1, +1\}$

Training Data (Vector Form)

ID	x_i	y
1	[0, 0, 0, 0]	-1
2	[0, 0, 0, 1]	-1
3	[0, 1, 0, 1]	+1
4	[0, 1, 1, 0]	+1
5	[0, 0, 0, 0]	-1
6	[1, 0, 0, 0]	-1
7	[1, 0, 0, 1]	-1
8	[1, 1, 1, 1]	+1
9	[1, 0, 1, 2]	+1
10	[1, 0, 1, 2]	+1
11	[2, 0, 1, 2]	+1
12	[2, 0, 1, 1]	+1
13	[2, 1, 0, 1]	+1
14	[2, 1, 0, 2]	+1
15	[2, 0, 0, 0]	-1

Each vector follows: $x_i = (\text{Age}, \text{Job}, \text{House}, \text{Credit})$

How the Weight Vector Is Influenced by the Data - Now, here's the part most people skip. SVM doesn't magically produce weights. It **learns them from patterns in the data**:

- Class labels
- Feature distributions
- And most importantly, how well each feature helps separate the classes

Analysis of the *Class : No* ($y = -1$)

We have 6 samples labeled as *No*. And something very clear shows up → Every single one of them has: Job = 0 and House = 0 → This is not a coincidence. This is a **pattern**.

It tells us: Not having a job or not owning a house (or both) → Strongly pushes a point toward the **negative side** of the boundary.

So from SVM's perspective: These features are **highly informative**.

Analysis of the *Class : Yes* ($y = +1$)

Now look at the positive class. We have 9 samples labeled Yes. Most of them follow: $Job = 1$, $House = 1$, and $Credit \geq 1$

There are a few exceptions, but they are rare. And here's an important idea: **SVM does not chase exceptions. It focuses on patterns that help maximize separation**

So overall, we see: $Job \rightarrow$ positive signal; $House \rightarrow$ positive signal; $Credit \rightarrow$ supportive signal

Feature Importance (SVM Perspective) - Now let's translate this into how SVM "thinks":

- *Job and House*
 - Provide clear separation
 - Consistent across classes
 - Strong impact on margin→ These will get **large weights**
- *CreditRating*
 - Helps refine decisions
 - Especially within positive samples→ Positive weight, but smaller than *Job* and *House*
- *Age* - Now here's the twist. *Age* is... everywhere.
 - *Young* appears in both classes
 - *Middle* appears in both classes
 - *Old* appears in both classes→ No clear pattern. No separation.

So from SVM's perspective: *Age* does not help push the classes apart. And that leads to a very important conclusion: $w_{Age} \approx 0$. Not because *Age* is useless in general. But because in **this dataset**, it doesn't help.

SVM does not choose features based on intuition, frequency, or human assumptions. It chooses based on one thing:

▮ **Does this feature help maximize the margin?**

If yes → it gets weight; If not → it fades away

Final Interpretation - From everything we observed, the weights become:

- $w_1 = 0$ (*Age*); $w_2 > 0$ (*Job*); $w_3 > 0$ (*House*); $w_4 > 0$ (*Credit*) → So we can approximate: $w = (0, 1, 1, 1)$

And this vector is saying something very clear:

- "*Job, House, and Credit* push you toward *Yes*."
- "*Age* doesn't really matter here."

Effect on the Dot Product - Now something elegant happens. Since $w_1 = 0$, *Age* disappears completely: $w \cdot x = \text{Job} + \text{House} + \text{Credit}$

SVM has **automatically removed a feature**. No manual feature selection. No preprocessing tricks. It simply ignored what didn't help.

Decision Function - The general decision rule is:

$$f(x) = w \cdot x + b$$

Substituting our weights: $f(x) = \text{Job} + \text{House} + \text{Credit} + b$

And now, everything depends on one final piece → The bias b

Which must satisfy the core constraint: $y_i(w \cdot x_i + b) \geq 1$

And this is where things become precise. Because choosing b means deciding exactly **where the boundary sits**.

Next, we'll compute something very concrete → The dot product values for all samples. And that's where the geometry of SVM finally becomes visible.

Dot Product Values for All Samples - Now we finally arrive at something concrete. Up until now, everything was about *understanding*. This is where we start **computing**.

- Since we already determined: $w = (0, 1, 1, 1)$
- The dot product becomes: $w \cdot x = \text{Job} + \text{House} + \text{Credit}$

Computed Dot Products

ID	x_i	$w \cdot x_i$	y
1	[0, 0, 0, 0]	0	-1
2	[0, 0, 0, 1]	1	-1
3	[0, 1, 0, 1]	2	+1
4	[0, 1, 1, 0]	2	+1
5	[0, 0, 0, 0]	0	-1
6	[1, 0, 0, 0]	0	-1
7	[1, 0, 0, 1]	1	-1
8	[1, 1, 1, 1]	3	+1
9	[1, 0, 1, 2]	3	+1
10	[1, 0, 1, 2]	3	+1

ID	x_i	$w \cdot x_i$	y
11	[2, 0, 1, 2]	3	+1
12	[2, 0, 1, 1]	2	+1
13	[2, 1, 0, 1]	2	+1
14	[2, 1, 0, 2]	3	+1
15	[2, 0, 0, 0]	0	-1

And now something very clean appears:

- Negative class → values: **0, 1**
- Positive class → values: **2, 3**

There is a **gap** between them. That gap is where SVM will place the boundary.

Choosing the Bias b : Now we position the hyperplane. And SVM does this with one strict condition: $y_i(w \cdot x_i + b) = 1 \rightarrow$ At least one positive and one negative point must lie **exactly on the margin**.

- **Negative Class ($y = -1$):** Constraint becomes: $w \cdot x_i + b \leq -1$
 - The largest negative projection: $\max(w \cdot x_{\text{neg}}) = 1$
 - So: $1 + b = -1 \Rightarrow b = -2$
- **Positive Class ($y = +1$):** Constraint becomes: $w \cdot x_i + b \geq 1$
 - The smallest positive projection: $\min(w \cdot x_{\text{pos}}) = 2$
 - So: $2 + b = 1 \Rightarrow b = -1$

Maximum-Margin Choice: Now here's the key idea. SVM does not pick either extreme. It chooses the bias that **centers the margin** between the two groups → That gives: $b = -1$

Final Decision Function: $f(x) = \text{Job} + \text{House} + \text{Credit} - 1$

Classification Rule

- $f(x) \geq 0 \rightarrow \text{Yes}(+1)$
- $f(x) < 0 \rightarrow \text{No}(-1)$

Notice something subtle: b is not chosen randomly. It is not optimized independently. It is **forced into position** by the margin constraints. SVM doesn't "guess" the boundary. It *locks it exactly between the closest opposing points* → That's why SVM is stable, robust, and geometrically precise

Identifying the Support Vectors

Now we find the most important points. Support vectors satisfy: $y_i(w \cdot x_i + b) = 1$

- **Negative Support Vectors ($y = -1$)** - Condition: $w \cdot x_i + b = -1 \rightarrow$ Using $b = -1$:

ID	Vector	Reason
1	[0, 0, 0, 0]	$0 - 1 = -1$
5	[0, 0, 0, 0]	$0 - 1 = -1$
6	[1, 0, 0, 0]	$0 - 1 = -1$
15	[2, 0, 0, 0]	$0 - 1 = -1$

- **Positive Support Vectors ($y = +1$)** - Condition: $w \cdot x_i + b = +1$

ID	Vector	Reason
3	[0, 1, 0, 1]	$2 - 1 = 1$
4	[0, 1, 1, 0]	$2 - 1 = 1$
12	[2, 0, 1, 1]	$2 - 1 = 1$
13	[2, 1, 0, 1]	$2 - 1 = 1$

Key Observation - Only these points matter. Everything else? Irrelevant to the boundary (as long as they stay outside the margin).

This is one of the most beautiful properties of SVM: It builds the entire model using only a **small subset of the data**

Constructing the Hyperplane - The decision boundary is: $w \cdot x + b = 0$

Substitute: $0 \cdot \text{Age} + \text{Job} + \text{House} + \text{Credit} - 1 = 0$

Simplifies to: $\text{Job} + \text{House} + \text{Credit} = 1$

Margin Hyperplanes

- Positive margin: $\text{Job} + \text{House} + \text{Credit} = 2$
- Negative margin: $\text{Job} + \text{House} + \text{Credit} = 0$

These three parallel boundaries define the **decision boundary**, the **margin region**, and the **support vectors** → **Final Classifier**: $\hat{y} = \text{sign}(\text{Job} + \text{House} + \text{Credit} - 1)$

- **Predicting a New Sample** - Let's test it with a **new Customer (ID 16)**: $x = (2, 1, 1, 1)$
 - **Step 1: Dot Product** → $w \cdot x = 0 \cdot 2 + 1 + 1 + 1 = 3$
 - **Step 2: Add Bias** → $f(x) = 3 - 1 = 2$
 - **Step 3: Classification** → $\hat{y} = \text{sign}(2) = +1$
- **Final Prediction**: *Yes* → *Credit Granted*

This entire process might look mechanical. But it's actually very elegant.

- Features decide direction → w
- Data decides position → b
- Support vectors define everything

- And classification becomes pure geometry

SVM doesn't memorize. It doesn't guess → It **draws the safest possible boundary and trusts the margin**.

3. Support Vector Machines with Python

We will use the following breast cancer dataset:

[breast_cancer_wisconsin.csv](#)

Support Vector Machines is straightforwardly implemented using the Scikit-Learn library:

```
# Support Vector Machine with Python
import pandas as pd
from sklearn.model_selection import train_test_split

# Read the dataframe
df = pd.read_csv('breast_cancer_wisconsin.csv')
target = 'diagnosis'
features = df.columns.drop(target)
# Label Encoding for the target
df['diagnosis'] = df['diagnosis'].map({"B": 0, "M": 1})
# 80% training and 20% test
train_df, test_df = train_test_split(df, test_size=0.2, random_state=42)
X_train = train_df[features]
X_test = test_df[features]
y_train = train_df[target]
y_test = test_df[target]
```

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

model = SVC(kernel='linear', C=1.0)
model.fit(X_train, y_train)
# Classification report
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.956140350877193

- Reduce dimensionality with PCA, and then train a Support Vector Machine model and obtain the hyperplane equation:

```

from sklearn.decomposition import PCA

# Plot the hyperplane with PCA
pca = PCA(n_components=2)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)
# Train SVM in PCA space
svm_pca = SVC(kernel='linear', C=1.0)
svm_pca.fit(X_train_pca, y_train)
# For a linear SVM:  $w_1 x_1 + w_2 x_2 + b = 0$ 
w = svm_pca.coef_[0]
b = svm_pca.intercept_[0]
print(f"Hyperplane equation:")
print(f"{w[0]:.3f} * PC1 + {w[1]:.3f} * PC2 + {b:.3f} = 0")

```

Hyperplane equation:

$0.006 * PC1 + -0.018 * PC2 + 0.418 = 0$

- Plot decision boundary, margins, and support vectors with PCA

```


import numpy as np
import matplotlib.pyplot as plt

# Create mesh grid
x_min, x_max = X_train_pca[:, 0].min() - 1, X_train_pca[:, 0].max() + 1
y_min, y_max = X_train_pca[:, 1].min() - 1, X_train_pca[:, 1].max() + 1
xx, yy = np.meshgrid(
    np.linspace(x_min, x_max, 500),
    np.linspace(y_min, y_max, 500))
# Decision function
Z = svm_pca.decision_function(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
plt.figure(figsize=(10, 6))
# Decision boundary and margins
plt.contour(xx, yy, Z, levels=[-1, 0, 1], linestyles=['--', '-', '--'])

# Scatter plot
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train,
            cmap='coolwarm', edgecolors='k', alpha=0.7)
# Support vectors
plt.scatter(svm_pca.support_vectors_[:, 0],
            svm_pca.support_vectors_[:, 1], s=120,
            facecolors='none', edgecolors='k',
            linewidths=2, label='Support Vectors')

```

```
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("Linear SVM Hyperplane (PCA Projection)")
plt.legend()
plt.show()
```



Now you might be wondering. Why did we suddenly bring in PCA? Are we trying to improve the model? Optimize performance? Reduce noise? → Not this time. This is purely for **visualization**.

Because here's the problem: our dataset lives in a high-dimensional space—far beyond what we can see. And a hyperplane in that space? Impossible to imagine directly. So what PCA does is something very simple, yet very powerful—it compresses all those dimensions down into just two principal components, while preserving as much of the data's structure as possible. Not perfectly, but enough for us to see.

So when we plot the SVM after applying PCA, we're not looking at the exact original boundary—we're looking at its **projection**. A shadow of the real hyperplane, flattened into 2D, just so our eyes can finally understand what the model has been doing all along.

4. Support Vector Machines: Kernel Methods

Now everything we've done so far assumes something very convenient. That the data can be separated with a straight line or a plane, or higher. But what if it can't? What if the data looks like this: curved, or twisted, or completely overlapping in 2D. Does SVM fail? No. This is where SVM becomes **dangerously powerful**.

The Core Idea: Instead of forcing a linear boundary, SVM changes the space itself. Not by bending the line— But by **lifting the data into a higher dimension**, where a straight line *does* work. And here's the trick: It never actually performs this transformation explicitly.

The Kernel Trick: At the heart of this idea is something called the **kernel trick**. Instead of computing the transformation $\phi(x)$, SVM directly computes:

$$K(x, y) = \phi(x) \cdot \phi(y)$$

- So rather than saying: "Let's move the data to a higher dimension"
- It says: "Let's behave *as if* we already did"

This avoids massive computations, high-dimensional storage, and unnecessary complexity.

Why This Works - SVM doesn't actually need the coordinates. It only needs the **dot products between points**. So if we can compute those in a transformed space, we get all the benefits without the cost.

Robustness of Kernels - Now here's something subtle—but important. Real-world data is messy: Noise; Outliers; Imperfect separation. A good kernel doesn't try to fit everything perfectly. It focuses on the **meaningful relationships between points**. That's what makes kernels **robust**. They don't overreact to small fluctuations. They don't chase every outlier. They preserve structure without overfitting.

Common Kernels - Let's look at the most important ones.

- **Linear Kernel** The simplest case: $K(x, y) = x \cdot y \rightarrow$ No transformation. Just a standard linear SVM.
- **Polynomial Kernel:** $K(x, y) = (x \cdot y + c)^d$

This allows curved boundaries by introducing interactions between features.

Think of it as $\rightarrow$ "Combine features... then separate"

- **Radial Basis Function (RBF) Kernel:** $K(x, y) = \exp(-\gamma \|x - y\|^2)$

This one is the most widely used. Why? Because it thinks differently. Instead of drawing a global boundary, it asks: *How close are two points?*

- Nearby points $\rightarrow$ strong influence
- Far points $\rightarrow$ almost irrelevant

This allows it to create **very flexible, non-linear boundaries**.

Choosing the Kernel - Now comes the real challenge. Choosing a kernel is not about memorization. It's about understanding the data.

- Linear $\rightarrow$ when separation is simple
- Polynomial $\rightarrow$ when interactions matter
- RBF $\rightarrow$ when structure is complex and unknown

And parameters (like γ) control how flexible the model becomes. Too flexible $\rightarrow$ overfitting. Too rigid $\rightarrow$ underfitting $\rightarrow$ So again, balance.

Kernel methods do something almost philosophical. They say: "*If the problem looks impossible... change the space*". And suddenly, what looked non-linear becomes linear. What looked complex becomes separable.

SVM does not try to chase complexity. It doesn't twist the boundary endlessly. Instead, it changes perspective. It moves the data into a world where separation becomes simple. And

then? It draws a straight line. That's the elegance of kernel methods. Not making the model more complicated. But making the problem easier to solve.

5. Margins in Support Vector Machine

Now we arrive at the *heart* of SVM. Not the hyperplane. Not the kernel. But something even more important → **The margin**.

In SVM, the margin is the distance between the decision boundary (the hyperplane) and the closest data points from each class. And those closest points? You already know them → **Support vectors**. They are not just "important". They are the ones that **define the entire model**.

The True Objective of SVM - SVM is not trying to just separate classes. It is trying to do something much more specific → **Maximize the margin**

Because a larger margin means more confidence, better generalization, and less sensitivity to noise.

So instead of asking: "Can I separate the data?", SVM asks:

┆ *"How far apart can I push these classes... while still separating them?"*

- **Hard Margin SVM** - Let's start with the ideal case. Perfect data. No noise. Clean separation. In this case, SVM enforces:

$$y_i(w \cdot x_i + b) \geq 1$$

Every point must be correctly classified and stay outside the margin → No exceptions.

Sounds perfect, right? It is. And that's exactly why it almost never works in reality. Because real-world data is messy → Noise exists; outliers exist; and overlap exists. And a single outlier can completely break a hard-margin model.

- **Soft Margin SVM**

So instead of forcing perfection, SVM relaxes. It allows some points inside the margin, and some misclassifications. This is done using **slack variables** ξ_i :

$$y_i(w \cdot x_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$$

Now the model says → *"Stay outside the margin... unless you really can't."*

This small change makes a huge difference. Because now SVM becomes flexible, robust, and practical.

- **The Role of C**

Now comes the balancing act. SVM introduces a parameter → C . And this parameter controls everything.

Large C → Strict model → Punishes errors heavily, smaller margin, but with a risk of overfitting

Small C → Relaxed model → Allows more violations, wider margin, and better generalization

→ So C is essentially asking: “Do you want to be correct... or resilient?”

- **Optimization View:** At its core, SVM is solving an optimization problem.

- **Hard Margin:** $\min_{w,b} \frac{1}{2} \|w\|^2$

→ Subject to: $y_i(w \cdot x_i + b) \geq 1$

- **Soft Margin:** $\min_{w,b,\xi_i} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$

→ Subject to: $y_i(w \cdot x_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$

And if you look closely. Everything comes back to one idea: **Maximize margin, while controlling error**

Summary

Support Vector Machines might sound sophisticated And definitely they are. But at their core? They are surprisingly grounded. They don't try to be perfect. They don't try to classify everything correctly

→ Instead, they draw a boundary. Carefully. Deliberately. The one that keeps the two worlds as far apart as possible.

And if a few points fall on the wrong side? That's okay. Because SVM understands something important:

▮ ***There is no such thing as a perfect model. Only a well-balanced one.***

References

<https://www.geeksforgeeks.org/machine-learning/support-vector-machine-algorithm/>

<https://www.ibm.com/think/topics/support-vector-machine>

<https://www.youtube.com/watch?v=efR1C6CvbmE>

https://www.youtube.com/watch?v=_YPScrckx28

Next Section:

 29. One-vs-One and One-vs-All: Multi-Class Classification