



7. Supervised & Unsupervised Learning

The two main types of Machine Learning. At least for now

Previous Section:

 [6. Overfitting and Underfitting](#)

Contents:

[1. Supervised Learning](#)

[2. Unsupervised Learning](#)

[Summary](#)

[References](#)

We have learned how to prepare your data. But what type of problem is the data?

This is the moment where things start to split into different paths. Because not all data is meant to be learned in the same way. Sometimes, your dataset already comes with answers. Other times, it's just raw information, waiting to be understood.

And that difference? It leads us to two of the most important ideas in Machine Learning:


| Supervised Learning and Unsupervised Learning.

- In Supervised Learning, the model learns with guidance. You don't just give it data—you also give it the correct answers. It's like teaching a student with labeled examples: "This is a cat," "This is a dog." Over time, the model

begins to understand the relationship between inputs and outputs, so when it sees something new, it can make predictions based on what it has learned.

- Unsupervised Learning, on the other hand, is a completely different experience. There are no answers provided. No labels. Just data. The model is left to explore, to observe, and to discover patterns on its own. It might group similar data points together, detect hidden structures, or uncover relationships that were not obvious at first glance.

Reference the image below:



So while supervised learning is about **learning from known answers**, unsupervised learning is about **finding meaning where no answers exist**. And understanding which type of problem you are dealing with... is often the first real step toward choosing the right model.



If I were to describe it in a simpler way:

- Supervised Learning feels like a student who is guided step by step. You are given formulas, shown solved examples, and taught what the correct answers look like. Over time, you don't just memorize—you begin to recognize patterns, and when a new problem appears, you apply what you've learned to reach a solution.
- Unsupervised Learning, however, is a much quieter journey. No formulas. No answers. Just problems placed in front of you. At first, it feels uncertain, but as you observe more, you start noticing patterns. Similar problems begin to group together, structures begin to form, and meaning slowly emerges—not because it was given to you, but because you discovered it yourself.

And in a way, that difference captures the essence of Machine Learning itself: Sometimes we learn from answers. And sometimes, we learn by finding them.

1. Supervised Learning

Supervised Learning is, perhaps the most intuitive place to begin.

If you remember how we talked about learning from experience, this is the version where the experience comes with **answers attached**. Every piece of data already knows what it is supposed to be. The model is not wandering in the dark—it is being guided, corrected, and gradually shaped by what is right and what is wrong.

You can think of it as learning through **observation with feedback**. The model sees an input, compares its prediction to the actual output, makes a mistake, and then adjusts itself. And this process repeats. Again and again... until the mistakes become smaller.

At its core, the workflow is surprisingly simple:

Input + Output → Train the Model → Predict on New Data

- You provide the model with data (inputs) and their correct answers (outputs)
- The model learns the relationship between them
- And once it has learned enough, it can face something new and make a prediction

But behind that simplicity lies something deeper: The model is not memorizing answers—it is trying to **understand the relationship** between input and output.

Types of Supervised Learning Models: Not all problems look the same. And because of that... supervised learning splits into two main directions.

- **Regression — Predicting Continuous Values:** Sometimes, the answer is not a category... but a number.
 - Predicting house prices
 - Estimating temperature
 - Forecasting sales

These are problems where the output exists on a continuous scale.

Common models include: Linear Regression; Polynomial Regression; Lasso Regression; and Ridge Regression

Here, the model is trying to capture **how variables change together**—how one thing influences another.

- **Classification — Predicting Categories:** Other times, the answer belongs to a group.
 - Spam or Not Spam
 - Fraud or Legitimate
 - Cat, Dog, Horse, or Bird

Now the goal is not to predict “how much”... but to decide **what it is**.

Common models include: Logistic Regression; Support Vector Machines (SVM); Decision Trees; Random Forest; and Naive Bayes.

And even within classification, there is a subtle distinction:

- Binary Classification → One input belongs to one of the two classes → Example: Spam or Not Spam, Yes or No

- **Multi-Class Classification** → One input belongs to one of the three classes or above → Example: Cat *or* Dog *or* Bird
- **Multi-Label Classification** → One input can belong to multiple classes at once → Example: An image that contains both a *cat* and a *dog*

That difference may seem small, but it completely changes how the model thinks.

Where is Supervised Learning Used? Because it relies on labeled data, supervised learning appears in many real-world systems:

- Image Classification → Recognizing objects in images
- Medical Diagnosis → Predicting diseases from patient data
- Fraud Detection → Identifying suspicious transactions
- Natural Language Processing (NLP) → Understanding and classifying text

In all of these cases, the key idea is the same:

We already know the answers for past data, and we want to predict the future.

Advantages: There's a reason supervised learning is so widely used—it works incredibly well when the data is prepared properly.

- It learns directly from labeled examples, making predictions more accurate
- The more high-quality data you provide, the better it becomes
- It is flexible enough to handle both:
 - Classification (categories)
 - Regression (numerical values)

In a sense, it is a very **structured way of learning**—clear inputs, clear outputs, clear goals.

Disadvantages: But that structure comes at a cost.

- It requires well-labeled data, and labeling data is expensive, time-consuming, and sometimes even subjective
- It can fall into the trap of **overfitting**—learning the training data too well, but failing on new data

- And because the world changes, the model must be updated regularly with new labeled data to remain useful

So while supervised learning is powerful, it is also dependent. Dependent on the quality of the data and the effort put into preparing it.

And if you look closely, you'll notice something important. Supervised Learning is not just about prediction. It is about **learning from what is already known to understand what is not yet seen.**

2. Unsupervised Learning

Unsupervised Learning is... where things become quiet. There are no answers here. No labels. No guidance telling the model what is right or wrong. Only data... and the expectation that something meaningful can still be found within it.

If supervised learning is about being taught... then unsupervised learning is about **observing until understanding emerges**. The model is not trying to predict a known outcome. Instead, it is trying to answer a different question: *"What structure exists in this data... even if no one told me?"*

At a glance, the process is even simpler than supervised learning:

Input Data → Model → Clusters or Patterns

- You provide raw, unlabeled data
- The model analyzes similarities, differences, and distributions
- And in the end... it organizes the data into meaningful structures

There is no "correct answer" to compare against. Only patterns, waiting to be discovered.

Types of Unsupervised Learning Models: Because there are no labels, unsupervised learning focuses on **understanding structure**. And it usually does this in two main ways.

- **Clustering — Finding Natural Groups:** Sometimes, the data naturally forms groups, we just don't see them at first.

Clustering algorithms try to answer: *"Which data points belong together?"*

Common methods include:

- K-Means Clustering

- Hierarchical Clustering
- DBSCAN (Density-Based Spatial Clustering of Applications with Noise)
- Gaussian Mixture Models (GMMs)

Each of these approaches grouping differently: Some rely on distance; Some build hierarchies; and Some focus on density

→ But the goal is the same: **Reveal hidden structure by grouping similar data together**

- **Dimensionality Reduction — Simplifying Complexity:** Other times, the problem is not grouping, but understanding. Modern datasets can have dozens, hundreds, or even thousands of features. Too complex to visualize. Too complex to interpret.

Dimensionality reduction methods try to answer: ***“Can we simplify this data... without losing what matters?”***

Common methods include:

- Principal Component Analysis (PCA)
- Singular Value Decomposition (SVD)
- Independent Component Analysis (ICA)

These techniques compress data into fewer dimensions while preserving important patterns. It's like turning a complex, high-dimensional space into something we can actually see and reason about.

Where is Unsupervised Learning Used?

Even without labels, unsupervised learning is incredibly powerful in real-world scenarios:

- Anomaly Detection → Finding unusual or suspicious behavior
- Scientific Discovery → Identifying unknown patterns in data
- Recommendation Systems → Discovering user preferences and similarities
- Customer Segmentation → Grouping customers based on behavior

In many cases, unsupervised learning is used as a **first step**—to explore the data before making decisions.

Advantages: There is a certain freedom in unsupervised learning.

- It does not require labeled data, which makes it much easier to apply in real-world situations
- It can simplify large and complex datasets while preserving important structures
- It is capable of discovering hidden patterns, relationships, and trends that humans might overlook

→ In a way, it allows the data to **speak for itself**.

Disadvantages: But that freedom comes with uncertainty.

- Without labeled outputs, it is difficult to measure accuracy or know if the results are “correct”
- The lack of guidance can lead to patterns that are unclear, unstable, or even misleading
- The groups formed often require human interpretation afterward to assign meaning
- It is sensitive to missing data, noise, and outliers, which can distort the discovered structure

So while unsupervised learning can reveal insights, it does not guarantee that those insights are meaningful without careful interpretation.

And this is what makes unsupervised learning so different. It does not learn from answers. It learns from **existence itself**. Patterns are not given. They are discovered.

Summary

Studying Machine Learning is not about jumping straight into algorithms. Before you even ask *how* a model works, you need to step back and understand two more important questions.

- First: **Where does the model learn from?**

That is your **data**—how it is collected, cleaned, and prepared. Because no matter how powerful a model is, if the data is flawed, the learning will be flawed too.

- Second: **What is the model trying to learn?**

Is it learning from answers (**Supervised Learning**), or discovering patterns on its own (**Unsupervised Learning**)?

This decision shapes everything that comes after. Only after understanding these two does the question of *how* begin to make sense.

- **How do we model the problem?**
- **How do we represent the relationships?**

→ That's where methods like Decision Trees, Boosting, or Clustering come in.

But if you start with *how* too early, you're not really learning Machine Learning. You're just choosing tools without understanding the problem.

References

<https://www.geeksforgeeks.org/machine-learning/supervised-unsupervised-learning/>

<https://www.ibm.com/think/topics/supervised-vs-unsupervised-learning>

<https://www.mygreatlearning.com/blog/supervised-vs-unsupervised-learning/>

<https://cloud.google.com/discover/supervised-vs-unsupervised-learning>

https://www.youtube.com/watch?v=aaKUJZqWZ_I

Next Section:

8. Linear Regression